

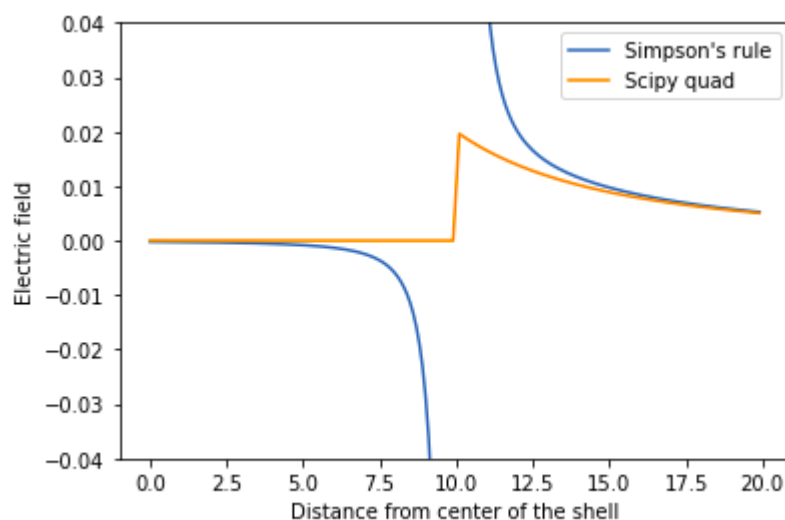
Q1. I take the formula for the electric field of the thin shell of charge density  $\sigma$  and radius  $R$  at a distance  $z$  from the centre of the shell from Griffiths solutions:

$$\frac{1}{4\pi\epsilon_0}(2\pi R^2\sigma) \int_{-1}^1 \frac{z - Ru}{(R^2 + z^2 - 2Rzu)^{3/2}} du$$

I take  $\sigma$  such that the factor outside the integral is 1. I take  $R = 10$  units, and  $z$  from 0 to 20 units. My code is based on Jon's code which uses the Simpsons rule of integration. I calculate the integration iteratively for different values of  $z$ . I also calculate the same integral using scipy quad for the different values of  $z$ :

```
R = 10
u = np.linspace(-1, 1, 100)
du = np.median(np.diff(u))
z = np.arange(0, 20, 0.1)
int_quad = np.zeros(len(z))
int_scipy = np.zeros(len(z))
for i in range(len(z)):
    fun = (z[i] - R*u)/(R**2 + z[i]**2 - 2*R*z[i]*u)**1.5
    int_quad[i] = du/3.0*(fun[0]+fun[-1]+4*np.sum(fun[1:-1:2])+2*np.sum(fun[2:-1:2]))
    scipy_fun = lambda x: (z[i] - R*x)/(R**2 + z[i]**2 - 2*R*z[i]*x)**1.5
    int_scipy[i], error = integrate.quad(scipy_fun, -1, 1)
```

The integral and thus the electric field varies like this with the distance from the centre of the shell:



As expected, the  $E$  is 0 for  $z < R$  and goes like  $1/z^2$  for  $z > R$ . There is a singularity at  $z = R$ , which my integration code cares about, but scipy quad ignores.

Q2. My function is based on Jon's function for adaptive step size integration. The idea is that we take 5 equally spaced points with step size  $dx$  between the integration limits  $a$  and  $b$ , and use the Simpson's rule to get a value of integration for it ( $i_2$ ). To estimate the error, we then take a step size of  $2dx$ , and thus use just the 1st, 3rd, and 5th of those equally spaced points (which are now separated by  $2dx$ ) to get the integration using Simpson's rule ( $i_1$ ). If  $i_2$  and  $i_1$

differ by more than the tolerance, we split the interval and two halves and call the function on them recursively with  $\text{tol}=\text{tol}/2$  until  $i2-i1<\text{tol}$ .

To avoid calling the function multiple times for the same  $x$ , I use the fact that the first call evaluates  $f(x)$  at  $a, a+dx, a+2dx, a+3dx, a+4dx=b$ . When the function is called recursively on the first half of the interval,  $dx$  will be halved and the function values will be needed at  $a, a+dx/2, a+dx, a+3dx/2, a+2dx$ . Since, the function is already evaluated at  $a, a+dx$  and  $a+2dx$  in the previous call, I pass that information while calling the function next. Similarly, I pass  $a+2dx, a+3dx, a+4dx$  while calling the function on the second half of the interval:

```
def integrate_adaptive(fun,a,b,tol,extra=None):
    fun_count = 0
    x=np.linspace(a,b,5)
    dx=x[1]-x[0]
    if extra == None:
        y=fun(x)
        fun_count += len(x)
    else :
        y = [extra[0], fun(x[1]), extra[1], fun(x[3]), extra[2]]
        fun_count += 2
    #do the 3-point integral
    i1=(y[0]+4*y[2]+y[4])/3*(2*dx)
    i2=(y[0]+4*y[1]+2*y[2]+4*y[3]+y[4])/3*dx
    myerr=np.abs(i1-i2)
    if myerr<tol:
        return i2,fun_count
    else:
        mid=(a+b)/2
        int1, count1=integrate_adaptive(fun,a,mid,tol/2, extra=[y[0], y[1], y[2]])
        int2, count2=integrate_adaptive(fun,mid,b,tol/2, extra=[y[2], y[3], y[4]])
        return int1+int2, count1+count2
```

As can be seen, my function has much lesser function call than the lazy approach we did in class:

1) For  $e^x$  between 0 and 2:

```
# Integrating e^x from 0 to 2
from scipy import integrate
ans1, count1 = integrate_adaptive(np.exp, 0, 2, tol = 1e-6, extra=None)
ans2, count2 = integrate_lazy(np.exp, 0, 2, tol = 1e-6)
print("Integral = {} and # of function calls = {} for my function".format(ans1, count1))
print("Integral = {} and # of function calls = {} for the lazy function".format(ans2, count2))
print("Real integral value = {}".format(integrate.quad(np.exp, 0, 2)[0]))

Integral = 6.3890561284651 and # of function calls = 34 for my function
Integral = 6.3890561284651 and # of function calls = 85 for the lazy function
Real integral value = 6.3890560989306495
```

2) For  $\ln(x)$  between 1 and 10:

```
ans1, count1 = integrate_adaptive(np.log, 1, 10, tol = 1e-6, extra=None)
ans2, count2 = integrate_lazy(np.log, 1, 10, tol = 1e-6)
print("Integral = {} and # of function calls = {} for my function".format(ans1, count1))
print("Integral = {} and # of function calls = {} for the lazy function".format(ans2, count2))
print("Real integral value = {}".format(integrate.quad(np.log, 1, 10)[0]))
```

Integral = 14.025850907644363 and # of function calls = 76 for my function  
 Integral = 14.025850907644363 and # of function calls = 190 for the lazy function  
 Real integral value = 14.025850929940457

3) For  $\sin(x)$  between 0 and  $\pi/2$

```
ans1, count1 = integrate_adaptive(np.sin, 0, np.pi, tol = 1e-6, extra=None)
ans2, count2 = integrate_lazy(np.sin, 0, np.pi, tol = 1e-6)
print("Integral = {} and # of function calls = {} for my function".format(ans1, count1))
print("Integral = {} and # of function calls = {} for the lazy function".format(ans2, count2))
print("Real integral value = {}".format(integrate.quad(np.sin, 0, np.pi)[0]))
```

Integral = 2.0000000217516316 and # of function calls = 48 for my function  
 Integral = 2.0000000217516316 and # of function calls = 120 for the lazy function  
 Real integral value = 2.0

Q3. I write a function which takes in  $x$  from 0.5 to 1, rescales it to get  $x$  from -1 to 1. Then, the function starts with fitting a Chebyshev polynomial of order 150 to  $y = \log_2(x)$ . Since we want an accuracy better than  $10^{-6}$ , the function iteratively reduces the order of the Chebyshev polynomial until the maximum error of the fit, that is, the maximum of  $y_{\text{predicted}} - y_{\text{real}}$  is just less than  $10^{-6}$ :

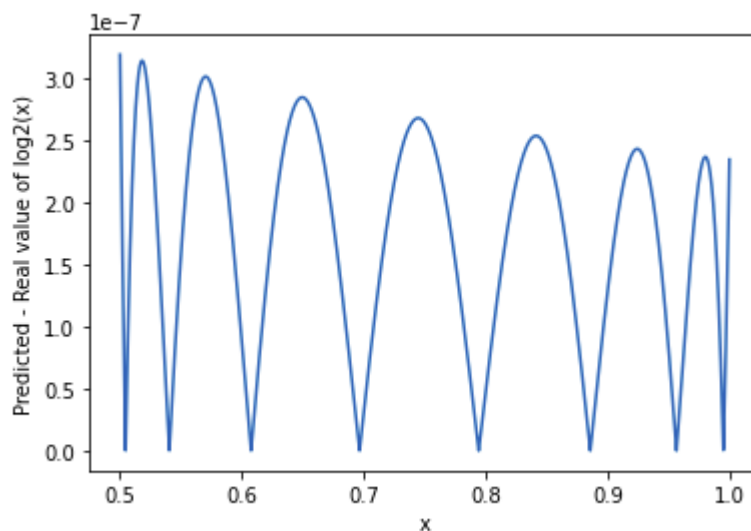
```
def cheby_model(x, y, tol, ord):
    len_x = len(x)
    x_min = x[0]
    x_max = x[len(x)-1]
    old_range = x_max - x_min
    new_range = 1 - (-1)
    new_x = (((x - x_min) * new_range) / old_range) + (-1)
    mat = np.zeros([len_x, ord+1])
    mat[:,0] = 1.0
    if ord > 0:
        mat[:,1] = new_x
    if ord > 1:
        for i in range(1, ord):
            mat[:,i+1] = 2 * new_x * mat[:,i] - mat[:,i-1]
    coeff = np.polynomial.chebyshev.chebfit(new_x, y, ord)
    pred = np.dot(mat, coeff)
    error = np.abs(y - pred)
    max_error = np.max(error)
    while(max_error >= tol):
        ord = ord - 1
        mat = mat[:, 0:ord+1:1]
        coeff = coeff[0:ord+1:1]
        pred = np.dot(mat, coeff)
        error = np.abs(y - pred)
        max_error = np.max(error)
        rms_error = np.sqrt(np.mean((pred - y)**2))
    return pred, max_error, rms_error, ord+1, coeff
```

As can be seen, the accuracy of the fit is better than  $10^{-6}$  when we take 8 terms in the Chebyshev polynomial (that is, a 7 order Chebyshev polynomial):

```
x = np.linspace(0.5, 1, 10000)
y = np.log2(x)
ord = 150
tol = pow(10, -7)
pred, max_error, rms_error, terms, coeff = cheby_model(x, y, tol, ord)
print("The number of terms needed to have a max error = {} and rms error = {} is {}".format(max_error, rms_error, terms))
```

The number of terms needed to have a max error = 3.196978304309539e-07 and rms error = 1.9185635134716262e-07 is 8

The residuals of the fit look like:



Then, I define a function to take the  $\ln(x)$ . For that, for each input  $x$ , I call `np.frexp(x)` which splits  $x$  as:  $x = \text{mant} \cdot 2^{\text{expo}}$ .  $\log_2$  of  $x$  is then just  $\log_2(\text{mant}) + \text{expo}$ . Mant is between 0.5 to 1, so the function defined above can be used to get  $\log_2(\text{mant})$ . I do the same procedure to get  $\log_2(e)$ . The  $\ln(x)$  is just  $\log_2(x)/\log_2(e)$ .

```
def mylog2(x, coeff):
    mant, exp = np.frexp(x)
    mant_new = (mant - 0.5)*(1 - (-1))/(1-0.5) + (-1)
    log2_x = np.polynomial.chebyshev.chebval(mant_new, coeff) + exp
    mant_e, exp_e = np.frexp(np.e)
    mant_e_new = (mant_e - 0.5)*(1 - (-1))/(1-0.5) + (-1)
    log2_e = np.polynomial.chebyshev.chebval(mant_e_new, coeff) + exp_e
    ln = log2_x/log2_e
    return ln
```

The residuals and errors for the  $\ln(x)$  calculated by the above function looks like:

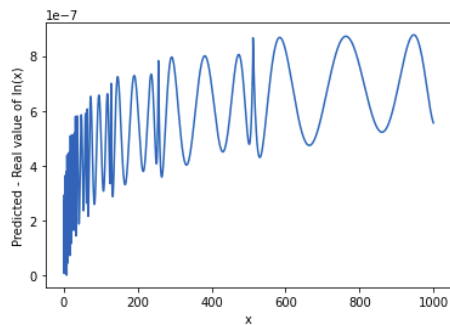
```

x = np.linspace(0.1, 1000, 10000)
ln_pred = mylog2(x, coeff)
ln = np.log(x)
max_error_ln = np.max(np.abs(ln_pred - ln))
rms_error_ln = np.sqrt(np.mean((ln_pred - ln)**2))
print("Max error = {} and RMS error = {} in the chebyshev fit for ln(x)".format(max_error_ln, rms_error_ln))
plt.plot(x, np.abs(ln_pred - ln))
plt.xlabel("x")
plt.ylabel("Predicted - Real value of ln(x)")

```

Max error = 8.768117671920095e-07 and RMS error = 6.324117839567934e-07 in the chebyshev fit for ln(x)

```
Text(0, 0.5, 'Predicted - Real value of ln(x)')
```



Repeating the same procedure as above for legendre polynomial fit of order 7 (order as returned by my cheby\_model function above):

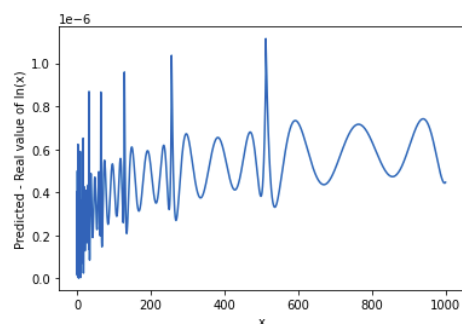
```

: x = np.linspace(0.1, 1000, 10000)
ln_pred_leg = mylog2_legendre(x, coeff_legendre)
ln = np.log(x)
max_error_ln_leg = np.max(np.abs(ln_pred_leg - ln))
rms_error_ln_leg = np.sqrt(np.mean((ln_pred_leg - ln)**2))
print("Max error = {} and RMS error = {} in the legendre fit for ln(x)".format(max_error_ln_leg, rms_error_ln_leg))
plt.plot(x, np.abs(ln_pred_leg - ln))
plt.xlabel("x")
plt.ylabel("Predicted - Real value of ln(x)")

```

Max error = 1.1141421474292201e-06 and RMS error = 5.426786559419055e-07 in the legendre fit for ln(x)

```
: Text(0, 0.5, 'Predicted - Real value of ln(x)')
```



As can be seen, the max error is lower for the Chebyshev fit, but the RMS error is lower for the Legendre fit.